

Concept Learning Data Analysis

Sahil Chopra

May 2, 2018

Introduction

Here, we present the preliminary analysis of the concept learning data collected April 30 - May 2, 2018. Before we delve into our analysis, we summarize the stimuli and concepts utilized in this experiment.

The stimuli have four axes of variability, each with three possible values – allowing for 81 unique critters. 50 critters were used in training and 31 were held out for test. Teachers and Listeners, who were paired together, were provided the same critters during the test set.

The four axes of variability were as follows:

- Critter Type (Bug, Fish, Bird)
- Primary Color (Blue, Green, Orange)
- Secondary Color (Red, Yellow, Purple)
- Size (Small, Medium, Large)

We intentionally omitted rules that relied on size. Unlike the other three properties, size is used described in relative terms, e.g. “small”, “medium”, “large”. Without visual grounding, it’s unclear whether a listener would immediately understand what these terms refer to, when presented the test set. We will address this issue in later iterations of this experiment.

We ran 5 concepts:

- 1 Single Feature Concept
- 2 Logical Conjunctions
- 2 Logical Disjunctions

For each concept, there were two lists. Each list comprised of the same stimuli, with different test / train splits and orderings of stimuli. We ran two lists to make sure that learning at similar rates was possible for a concept irrespective of the specific ordering of stimuli.

The 5 Concepts were:

- Primary Color == Orange
- Critter Type == Fish && Primary Color == Blue
- Primary Color == Orange && Secondary Color == Purple
- Critter Type == Bug || Secondary Color == Yellow
- Critter Type == Bird || Primary Color == Green

Data Processing

```
library(reshape2)
library(purrr)
library(jsonlite)
```

```

##
## Attaching package: 'jsonlite'

## The following object is masked from 'package:purrr':
##
##     flatten
library(tidyr)

##
## Attaching package: 'tidyr'

## The following object is masked from 'package:reshape2':
##
##     smiths
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##     filter, lag

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
library(ggplot2)
library(scales)

##
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
##
##     discard
# Load Training Trial Responses
temp <- list.files(
  "../mturk/mp-game-3/experiment_1/results-cleaned/train_trials",
  pattern="*.csv",
  full.names=TRUE
)
train_trials <- do.call(rbind, lapply(temp, read.csv))

# Load Test Trial Responses
temp <- list.files(
  "../mturk/mp-game-3/experiment_1/results-cleaned/test_trials",
  pattern="*.csv",
  full.names=TRUE
)
test_trials <- do.call(rbind, lapply(temp, read.csv))

# Load Training Summary Stats
temp <- list.files(
  "../mturk/mp-game-3/experiment_1/results-cleaned/train_summary_stats",
  pattern="*.csv",

```

```

    full.names=TRUE
  )
  train_stats <- do.call(rbind, lapply(temp, read.csv))

  # Load Test Summary Stats
  temp <- list.files(
    "../..//mturk/mp-game-3/experiment_1/results-cleaned/test_summary_stats",
    pattern="*.csv",
    full.names=TRUE
  )
  test_stats <- do.call(rbind, lapply(temp, read.csv))


```

Dataset Composition

```

rule_freq <- as.data.frame(table(train_stats$rule_idx)) %>%
  rename(rule_idx = Var1) %>%
  rename(num_games = Freq)
ggplot(rule_freq, aes(x=rule_idx, y = num_games)) + geom_bar(stat="identity")

```



Due to some server issues early on we have a slight imbalance in the number of trials prescribed to each rule index. We intended to have ~3 per rule type originally.

Analysis: Accuracy

First, we examine how people perform on the train and test splits, for each specific list.

```
num_train = 5
num_test = 31


train_acc <- train_stats %>%
  mutate(correct = hits + correct_rejections) %>%
  mutate(incorrect = misses + false_alarms) %>%
  select(game_id, rule_idx, correct, incorrect)

test_acc_teacher <- test_stats %>%
  filter(role == "explorer") %>%
  mutate(correct = hits + correct_rejections) %>%
  mutate(incorrect = misses + false_alarms) %>%
  select(game_id, rule_idx, correct, incorrect)

test_acc_listener <- test_stats %>%
  filter(role == "student") %>%
  mutate(correct = hits + correct_rejections) %>%
  mutate(incorrect = misses + false_alarms) %>%
  select(game_id, rule_idx, incorrect, correct)
```


Training

```
temp <- melt(train_acc, id.var= c("game_id", "rule_idx"))
temp$variable <- factor(temp$variable, levels = c("incorrect", "correct"))
training_acc_plot <- ggplot(temp, aes(x=game_id, y=value, fill=variable)) +
  geom_bar(stat = "identity") +
  facet_grid(.~rule_idx, scales="free_x", space="free") +
  ylab("Responses") +
  theme(axis.text.x=element_text(angle=90)) +
  scale_x_discrete(label=abbreviate) +
  labs(title = "Train Accuracy (Teacher)") +
  scale_fill_manual( values = c("#3ABFC3", "#F8766D"))
plot(training_acc_plot)
```




Test

```
temp <- melt(test_acc_teacher, id.var= c("game_id", 'rule_idx'))
temp$variable <- factor(temp$variable, levels = c("incorrect", "correct"))
test_acc_teacher_plot <- ggplot(temp, aes(x=game_id, y=value, fill=variable)) +
  geom_bar(stat = "identity") +
  facet_grid(.~rule_idx, scales="free_x", space="free") +
  ylab("Responses") +
  theme(axis.text.x=element_text(angle=90)) +
  scale_x_discrete(label=abbreviate) +
  labs(title = "Test Accuracy (Teacher)") +
  scale_fill_manual( values = c("#3ABFC3", "#F8766D"))
plot(test_acc_teacher_plot)
```



```
temp <- melt(test_acc_listener, id.var= c("game_id", 'rule_idx'))
temp$variable <- factor(temp$variable, levels = c("incorrect", "correct"))
test_acc_listener_plot <- ggplot(temp, aes(x=game_id, y=value, fill=variable)) +
  geom_bar(stat = "identity") +
  facet_grid(.~rule_idx, scales="free_x", space="free") +
  ylab("Responses") +
  theme(axis.text.x=element_text(angle=90)) +
  scale_x_discrete(label=abbreviate) +
  labs(title = "Test Accuracy (Listener)") +
  scale_fill_manual( values = c("#3ABFC3", "#F8766D"))
plot(test_acc_listener_plot)
```



Note that test accuracy of listeners is often very similar to the test accuracy of teacher. There are few notable exceptions.

For rule_idx 3 game 1264, we see that the student performs much worse than the teacher. Upon examination of the chat logs, it seems like the student “hit” continue without conversing with the explorer. They only said “hello” to each other, before moving on to the test set. Thus, their performance is indicative of random guessing.

We also see significant discrepancies in performance on rule_idx 5. The rule here was Primary Color == Orange && Secondary Color == Purple. In Game 34589, the teacher correctly understood the concept, but said they also thought that fish couldn’t be wudsy, mentioning that he had never seen orange and purple fish during training. In Game 9542, the teacher told the student that “I didn’t notice any patterns at all.” – so again the student was left with no information, i.e. an uninformative prior.

Analysis: Hits/Misses and Correct Rejections/False Alarms

Nexr, we break down how people perform on the train and test splits, for each specific list.

```
num_train = 5
num_test = 31

train_perf <- train_stats %>%
  select(game_id, rule_idx, hits, misses, correct_rejections, false_alarms)


test_perf_teacher <- test_stats %>%
  filter(role == "explorer") %>%
  select(game_id, rule_idx, hits, misses, correct_rejections, false_alarms)
```

```

test_perf_listener <- test_stats %>%
  filter(role == "student") %>%
  select(game_id, rule_idx, hits, misses, correct_rejections, false_alarms)

temp <- melt(train_perf, id.var= c("game_id", 'rule_idx'))
temp$variable <- factor(temp$variable, levels = c("hits", "misses", "correct_rejections", "false_alarms"))
training_perf_plot <- ggplot(temp, aes(x=game_id, y=value, fill=variable)) +
  geom_bar(stat = "identity") +
  facet_grid(.~rule_idx, scales="free_x", space="free") +
  ylab("Responses") +
  theme(axis.text.x=element_text(angle=90)) +
  scale_x_discrete(label=abbreviate) +
  labs(title = "Train Performance (Teacher)")
plot(training_perf_plot)

```



Note that the bars aren't always split evenly across games with the same rule – except rule 7. I have to investigate why that's the case, as this shouldn't happen. We don't see this below for the test data.


```

temp <- melt(test_perf_teacher, id.var= c("game_id", 'rule_idx'))
temp$variable <- factor(temp$variable, levels = c("hits", "misses", "correct_rejections", "false_alarms"))
test_perf_teacher_plot <- ggplot(temp, aes(x=game_id, y=value, fill=variable)) +
  geom_bar(stat = "identity") +
  facet_grid(.~rule_idx, scales="free_x", space="free") +
  ylab("Responses") +
  theme(axis.text.x=element_text(angle=90)) +
  scale_x_discrete(label=abbreviate) +
  labs(title = "Test Performance (Teacher)")

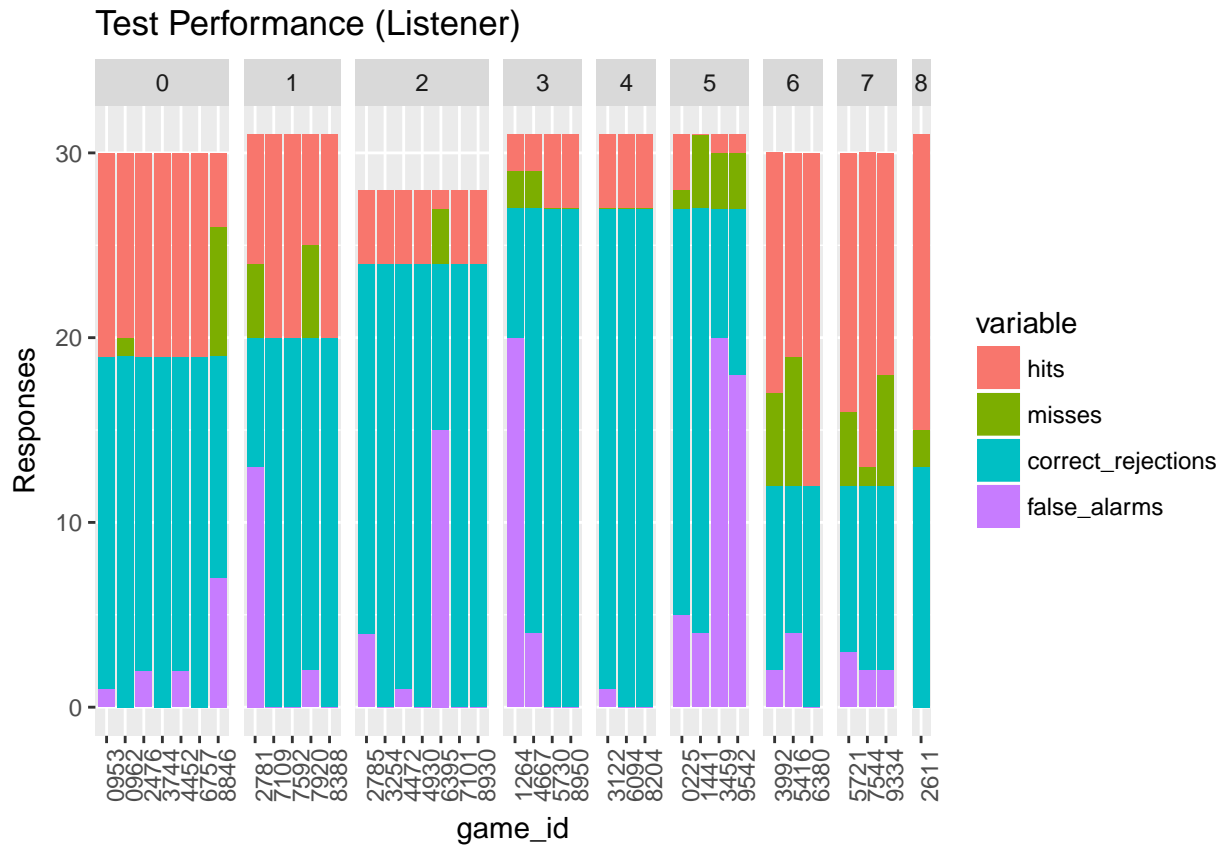
```



```
plot(test_perf_teacher_plot)
```



```
temp <- melt(test_perf_listener, id.var= c("game_id", 'rule_idx'))
temp$variable <- factor(temp$variable, levels = c("hits", "misses", "correct_rejections", "false_alarms"))
test_perf_listener_plot <- ggplot(temp, aes(x=game_id, y=value, fill=variable)) +
  geom_bar(stat = "identity") +
  facet_grid(.~rule_idx, scales="free_x", space="free") +
  ylab("Responses") +
  theme(axis.text.x=element_text(angle=90)) +
  scale_x_discrete(label=abbreviate) +
  labs(title = "Test Performance (Listener)")
plot(test_perf_listener_plot)
```



Performance for the listeners generally tracked the performance of the teachers on the test set. There are few exemplars that stand out though. Game 1264 in rule 3 had no dialogue, as discussed earlier. Games 3459 and 9452 have uninformative or poorly interpreted dialogue, as discussed earlier.

Analysis: Rational Rules

Posterior Predictives

Rule Selection

Analysis: Hold One Out Predictions